

Amazon DynamoDB

What is NoSQL

- NoSQL databases are non-tabular databases and store data differently than relational tables
- NoSQL databases store data in documents (ex: json) rather than relational tables
- NoSQL databases do not support join
- NoSQL databases scale horizontally
- Benefits of NoSQL Databases
 - Flexible data models
 - Horizontal scaling
 - Faster queries
 - Easy to work

NoSQL Database

- Most-popular types of NoSQL database
 - Document databases
 - Key-value stores
 - Wide-column databases
 - Graph databases
- Popular NoSQL Databases
 - MongoDB
 - Couchbase
 - Redis
 - Amazon DynamoDB
 - IBM Cloudant
 - RavenDB
 - Cassandra
 - HBase
 - Azure Cosmos DB

Amazon DynamoDB

- Amazon DynamoDB is a fully managed, serverless, key-value NoSQL database designed to run high-performance applications at any scale.
- Single-digit millisecond performance at any scale
- Millions of requests per seconds, trillions of row, 100s of TB of storage
- Each table has a **primary key**
- Each table can have an infinite number of items
- Maximum size of a item is 400KB
- Data types supported are:
 - Scalar Types: String, Number, Binary, Boolean, Null
 - Document Types: List, Map
 - Set Types: String Set, Number Set, Binary Set

Components of Amazon DynamoDB

- The core components of DynamoDB are **tables**, **items**, and **attributes**
- A table is a collection of items, and each item is a collection of attributes.
- DynamoDB uses primary keys to uniquely identify each item in a table and secondary indexes to provide more querying flexibility.
- **Tables:** A table is a collection of data. DynamoDB stores data in tables.
- **Items:** Each table contains zero or more items. An item is a group of attributes that is uniquely identifiable among all of the other items.
- **Attributes:** Each item is composed of one or more attributes. An attribute is a fundamental data element, something that does not need to be broken down any further.

People

```
{  
  "PersonID": 101,  
  "LastName": "Smith",  
  "FirstName": "Fred",  
  "Phone": "555-4321"  
}
```

```
{  
  "PersonID": 102,  
  "LastName": "Jones",  
  "FirstName": "Mary",  
  "Address": {  
    "Street": "123 Main",  
    "City": "Anytown",  
    "State": "OH",  
    "ZIPCode": 12345  
  }  
}
```

```
{  
  "PersonID": 103,  
  "LastName": "Stephens",  
  "FirstName": "Howard",  
  "Address": {  
    "Street": "123 Main",  
    "City": "London",  
    "PostalCode": "ER3 5K8"  
  },  
  "FavoriteColor": "Blue"  
}
```

Music

```
{  
  "Artist": "No One You Know",  
  "SongTitle": "My Dog Spot",  
  "AlbumTitle": "Hey Now",  
  "Price": 1.98,  
  "Genre": "Country",  
  "CriticRating": 8.4  
}
```

```
{  
  "Artist": "No One You Know",  
  "SongTitle": "Somewhere Down The Road",  
  "AlbumTitle": "Somewhat Famous",  
  "Genre": "Country",  
  "CriticRating": 8.4,  
  "Year": 1984  
}
```

```
{  
  "Artist": "The Acme Band",  
  "SongTitle": "Still in Love",  
  "AlbumTitle": "The Buck Starts Here",  
  "Price": 2.47,  
  "Genre": "Rock",  
  "PromotionInfo": {  
    "RadioStationsPlaying": [  
      "KHCR",  
      "KQBX",  
      "WTNR",  
      "WJJH"  
    ]  
  },  
  "TourDates": {  
    "Seattle": "20150625",  
    "Cleveland": "20150630"  
  },  
  "Rotation": "Heavy"  
}
```

```
{  
  "Artist": "The Acme Band",  
  "SongTitle": "Look Out, World",  
  "AlbumTitle": "The Buck Starts Here",  
  "Price": 0.99,  
  "Genre": "Rock"  
}
```

DynamoDB – Primary Keys

- When you create a table, table name and the primary key will be given.
- The primary key uniquely identifies each item in the table, so that no two items can have the same key
- DynamoDB supports 2 types of primary keys
 - Partition key
 - A simple primary key, composed of one attribute known as the partition key
 - Partition key and sort key
 - A composite primary key, composed of two attributes. partition key and sort key.
 - The combination must be unique

DynamoDB – Provisioned Throughput

- Table must have provisioned read and write capacity units
- Read Capacity Units (RCU): throughput for reads
- Write Capacity Units (WCU): throughput for writes
- Option to setup auto-scaling of throughput to meet demand
- ReadCapacityUnits
 - The maximum number of strongly consistent reads consumed per second
- WriteCapacityUnits
 - The maximum number of writes consumed per second

DynamoDB – Write Capacity Units

- One write capacity unit represents one write per second for an item up to 1 KB in size.
- If the items are larger than 1 KB, more WCU are consumed
- **Example 1:** we write 10 objects per seconds of 2 KB each.
We need $2 * 10 = 20$ WCU
- **Example 2:** we write 6 objects per second of 4.5 KB each
We need $6 * 5 = 30$ WCU (4.5 gets rounded to the upper KB)
- **Example 3:** we write 120 objects per minute of 2 KB each
We need $120 / 60 * 2 = 4$ WCU

DynamoDB – Read Capacity Units

- One read capacity unit represents **one strongly consistent read** per second, or **two eventually consistent** reads per second, for an item up to 4 KB in size.
- If the items are larger than 4 KB, more RCU are consumed
- **Example 1:** 10 **strongly consistent** reads per seconds of 4 KB each
We need $10 * 4 \text{ KB} / 4 \text{ KB} = 10 \text{ RCU}$
- **Example 2:** 16 **eventually consistent** reads per seconds of 12 KB each
We need $(16 / 2) * (12 / 4) = 24 \text{ RCU}$
- **Example 3:** 10 strongly consistent reads per seconds of 6 KB each
We need $10 * 8 \text{ KB} / 4 = 20 \text{ RCU}$ (we have to round up 6 KB to 8 KB)

DynamoDB – Write, Delete Data

- PutItem - Write data to DynamoDB
- UpdateItem – Update data in DynamoDB
- DeleteItem
 - Delete an individual row
 - Ability to perform a conditional delete
- DeleteTable
 - Delete a whole table and all its items
 - Much quicker deletion than calling DeleteItem on all items

DynamoDB – Batching Writes

- BatchWriteItem
 - Up to 25 PutItem and / or DeleteItem in one call
 - Up to 16 MB of data written
 - Up to 400 KB of data per item
- Batching allows you to save in latency by reducing the number of API calls done against DynamoDB
- Operations are done in parallel for better efficiency

DynamoDB – Reading Data

- **GetItem:**
 - Read based on Primary key
 - Primary Key = HASH or HASH-RANGE
 - Eventually consistent read by default
 - Option to use strongly consistent reads (more RCU - might take longer)
 - ProjectionExpression can be specified to include only certain attributes
- **BatchGetItem:**
 - Up to 100 items
 - Up to 16 MB of data
 - Items are retrieved in parallel to minimize latency

Amazon DynamoDB Accelerator (DAX)

- It is a fully managed, highly available, in-memory cache for Amazon DynamoDB that delivers up to a **10 times performance** improvement—**from milliseconds to microseconds**—even at millions of requests per second.
- DAX does all the heavy lifting required to add in-memory acceleration to your DynamoDB tables, without requiring developers to manage cache invalidation, data population, or cluster management.