

Amazon Elastic Beanstalk

Problems on AWS for a Developer

- Managing/Creation Infrastructure
- Code Deployment
- Configuring databases
- Configuring load balancers
- Auto Scaling
- Monitoring
- Logging

Amazon Elastic Beanstalk

- AWS Elastic Beanstalk is the fastest and simplest way to get an application up and running on AWS.
- Upload your application code, and the service automatically handles all of the details, such as
 - Resource provisioning
 - Load balancing
 - Auto Scaling
 - Monitoring.
- You can quickly deploy and manage applications on the AWS cloud
- Reduces management complexity

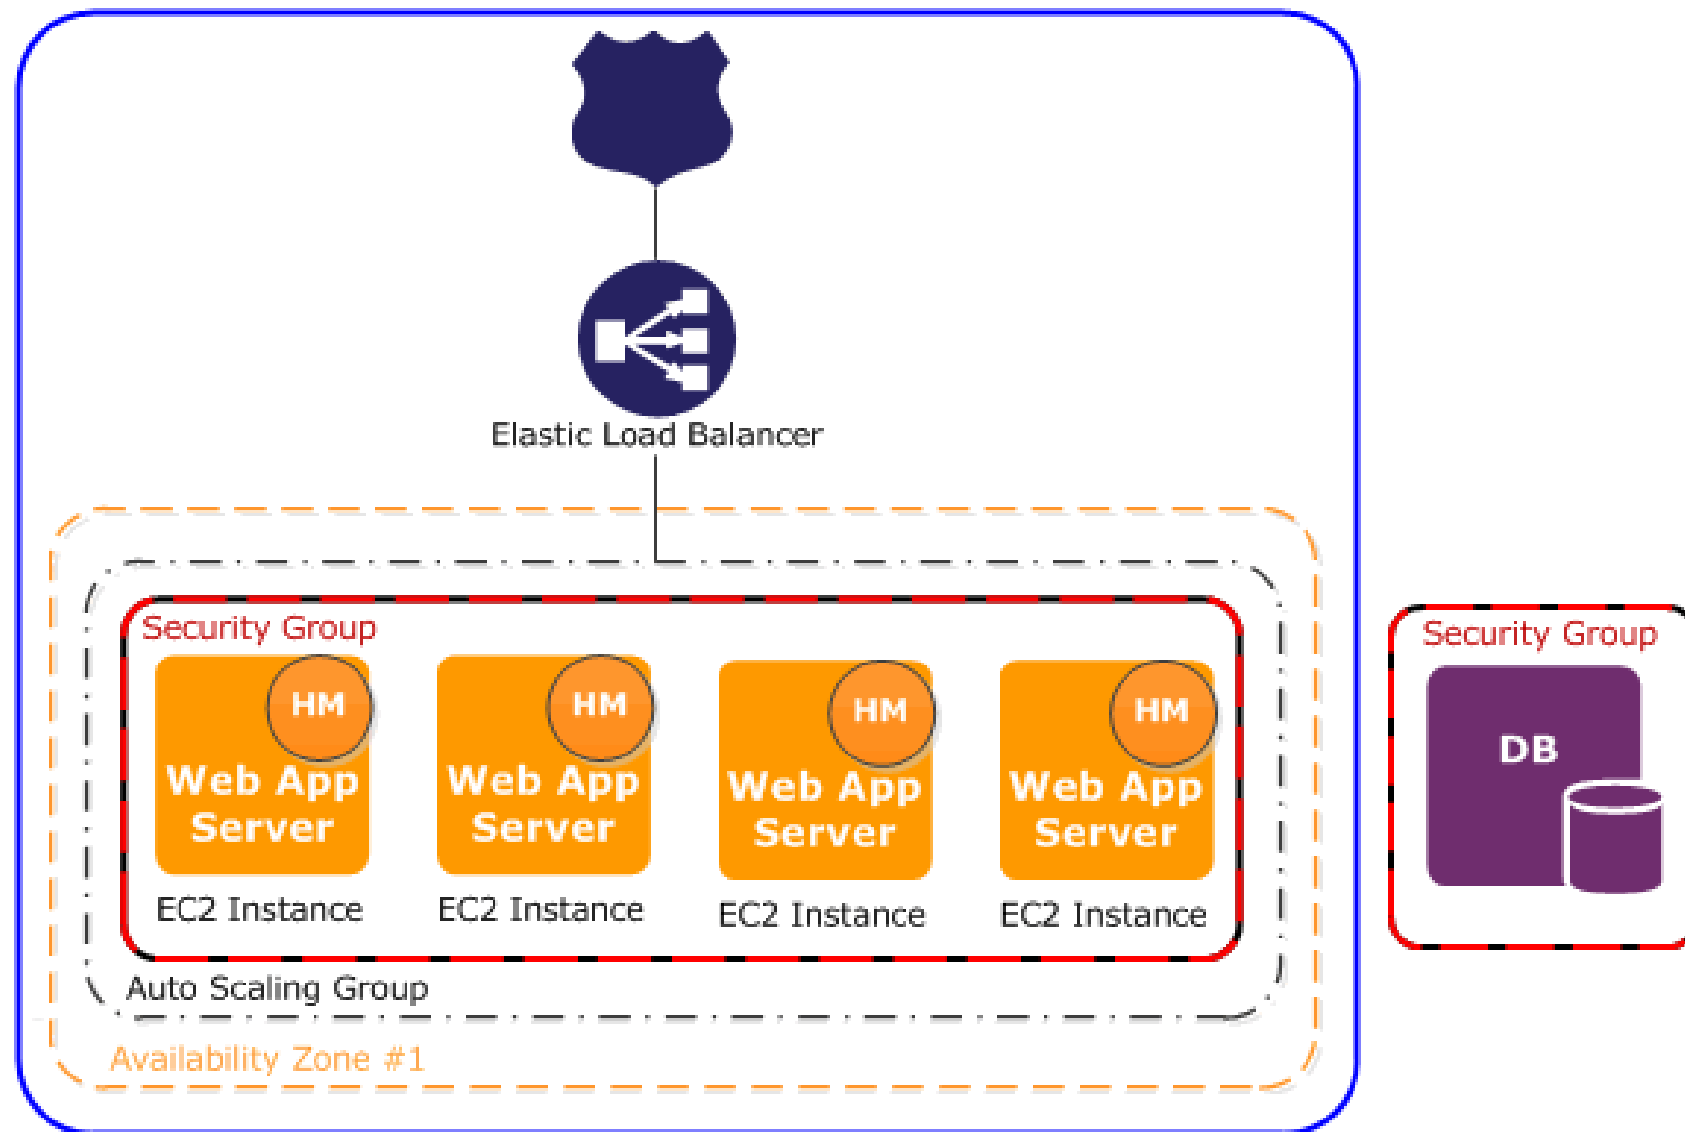
AWS Elastic Beanstalk Overview

- Fully Managed Service
 - Instance configuration / OS is handled by Beanstalk
 - Deployment strategy is configurable but performed by Elastic Beanstalk
- Just the application code is the responsibility of the developer
- Three architecture models:
 - Single Instance deployment: good for dev
 - LB + ASG: great for production or pre-production web applications
 - ASG only: great for non-web apps in production (workers, etc..)

AWS Elastic Beanstalk

- Elastic Beanstalk has three components
 - Application
 - Application version: each deployment gets assigned a version
 - Environment name (dev, test, prod...): free naming
- You deploy application versions to environments and can promote application versions to the next environment
- Rollback feature to previous application version
- Full control over lifecycle of environments

MyApp.elasticbeanstalk.com



Elastic Beanstalk

- Support for many platforms:
 - Go
 - .NET
 - Node.js
 - PHP
 - Python
 - Ruby
 - Packer Builder
 - Docker

Elastic Beanstalk Deployment

- Elastic Beanstalk supports environments as
 - Single Instance environments
 - High Availability with Load Balancer

Beanstalk Deployment Options

- **All at once**
 - Deployment fastest, but having little downtime
- **Rolling**
 - Deploy the new version in batches.
 - Each batch is taken out of service during the deployment phase, reducing your environment's capacity by the number of instances in a batch.
- **Rolling with additional batch:**
 - Deploy the new version in batches
 - first launch a new batch of instances to ensure full capacity during the deployment process
- **Immutable:**
 - Deploy the new version to a fresh group of instances
 - swaps all the instances when everything is healthy

Deployment Options - All at once

- Fastest deployment
- Application has a little downtime
- Great for quick iterations in development environment
- No additional cost
- rollback would take time in case of any issues

Deployment Options - Rolling


- Application is running below capacity
- Can set the bucket size
- Application is running both versions simultaneously
- No additional cost
- Long deployment

Deployment Options - Rolling with additional batch

- Application is running at capacity
- Can set the bucket size
- Application is running both versions simultaneously
- Small additional cost
- Additional batch is removed at the end of the deployment
- Longer deployment
- Good for prod env

Deployment Options - Immutable

- Zero downtime
- New Code is deployed to new instances on a temporary ASG
- High cost, double capacity
- Longest deployment
- Quick rollback in case of failures (just terminate new ASG)
- Good for prod env

| Method | Impact of failed deployment  | Deploy time | Zero downtime | No DNS change | Rollback process | Code deployed to |
|----------------------------------|---|-------------|---------------|---------------|---|----------------------------|
| All at once | Downtime | | No | Yes | Manual redeploy | Existing instances |
| Rolling | Single batch out of service; any successful batches before failure running new application version | † | Yes | Yes | Manual redeploy | Existing instances |
| Rolling with an additional batch | Minimal if first batch fails; otherwise, similar to Rolling | † | Yes | Yes | Manual redeploy | New and existing instances |
| Immutable | Minimal | | Yes | Yes | Terminate new instances | New instances |
| Traffic splitting | Percentage of client traffic routed to new version temporarily impacted | †† | Yes | Yes | Reroute traffic and terminate new instances | New instances |
| Blue/green | Minimal | | Yes | No | Swap URL | New instances |

Beanstalk Lifecycle Policy

- Elastic Beanstalk can store at most 1000 application versions
- If you don't remove old versions, you won't be able to deploy anymore
- To phase out old application versions, use a lifecycle policy
 - Based on time (old versions are removed)
 - Based on space (when you have too many versions)
- Versions that are currently used won't be deleted
- Option not to delete the source bundle in S3 to prevent data loss

Elastic Beanstalk Extensions

- A zip file containing our code must be deployed to Elastic Beanstalk
- All the parameters set in the UI can be configured with code using files
- Requirements:
 - in the `.ebextensions/` directory in the root of source code
 - YAML / JSON format
 - `.config` extensions (example: `logging.config`)
 - Ability to add resources such as RDS, ElastiCache, DynamoDB, etc...
- Resources managed by `.ebextensions` get deleted if the environment goes away